Boost Security in Kubernetes with CIS Security Controls and Benchmarks

Ader Fu <ader.ydfu@gmail.com> KCD Taipei 2024, Aug. 4, 2024



About me









Name:

Ader Fu, a.k.a. Yao-De Fu has extensive experience in DevOps.

Co-organizer:

- Cloud Native Taiwan User Group

 (https://www.facebook.com/groups/cloudnative.tw)
- Kubernetes Community Day(KCD Taiwan 2022&2023, KCD Taipei 2024)

Organization:

- Member: kubernetes, kubernetes-sigs, istio
- Teams: Kubernetes Special Interest Group (SIG) Docs

Community role:

- cncf/glossary approver(https://github.com/cncf/glossary)
- kubernetes/website approver and reviewer(https://github.com/kubernetes/website)

Certificate:

- 1st CNCF Kubestronauts in Taiwan (https://www.cncf.io/training/kubestronaut/)
- CKA, CKAD, CKS, KCNA and KCSA
- Prometheus Certified Associate (PCA)



kubestronaut



Outline



- Center for Internet Security (CIS) introduction
- Applying CIS benchmarks in Kubernetes
- Leveraging CNCF Landscape's security and compliance projects



Center for Internet Security (CIS)



What is system hardening?



A process of fortifying system configuration settings to safeguard against attack vectors—methods of gaining unauthorized access to computer or network systems—thereby reducing an organization's attack surface, which encompasses all possible points where an unauthorized user can access a system and extract data, to mitigate the risk of cyber-attacks.

Advantage:

- Reduce the risk of being attacked
- Improve system stability
- Meet compliance requirements



Configuration compliance scan



Configuration compliance scan is a Security Configuration Assessment (SCA) to verify if an asset's (Servers, Networking devices, IOT, Database, etc.) present configuration is based on best practice standards.

- Known System Hardening Configuration Standards
- Center for Internet Security (CIS) CIS
 Benchmarks(https://www.cisecurity.org/cis-benchmarks-overview)
- 2. National Institute of Standards and Technology (NIST) SP 800-123(General Server Security) and 800-53B(Security and Privacy Control)

(https://nvlpubs.nist.gov/nistpubs/legacy/sp/nistspecialpublication800-123.pdf) & (https://nvlpubs.nist.gov/nistpubs/SpecialPublications/NIST.SP.80053r5.pdf)

3. European Network and Information Security Agency (ENISA) (https://www.enisa.europa.eu)







CIS introduction





The Center for Internet Security, Inc. (CIS®)

- CIS are a community-driven nonprofit, responsible for the CIS Controls[®] and CIS Benchmarks[™], globally recognized best practices for securing IT systems and data.
- MS-ISAC and EI-ISAC are the key resources for cyber threat prevention, protection, response, and recovery for the nation's state, local, tribal, and territorial (SLTT) governments and elections communities, respectively



The CIS Vision

Leading the global community to secure our everchanging connected world.



The CIS Mission

Our mission is to make the connected world a safer place by developing, validating, and promoting timely best practice solutions that help people, businesses, and governments protect themselves against pervasive cyber threats.

What Types of IT Systems Do CIS Benchmarks Cover?

Cloud Providers

Alibaba Cloud
Amazon Web
Services
Google Cloud
Computing Platform
Microsoft 365
Microsoft Azure...

Desktop Software

Google Chrome
Microsoft Office
Microsoft Web
Browser
Mozilla Firefox
Safari Browser
Zoom...

DevSecOps Tools

Software Supply Chain Security



Mobiles Devices

Apple iOS Google Android



Multi Function Print Devices

Print Devices



Network Devices

Check Point Firewall

Cisco

F5

Fortinet

Juniper

Palo Alto Networks pfSense Firewall...

Operating System

Apple macOS

Azure Linux Bottlerocket

CentOS Linux

Debian Family Linux

Microsoft Windows

Server...

Server Software

Apache HTTP Server

Apache Tomcat

BIND

Docker

Kubernetes

MariaDB

Microsoft IIS...

Critical Security Controls CIS Controls





What are the CIS Benchmarks?

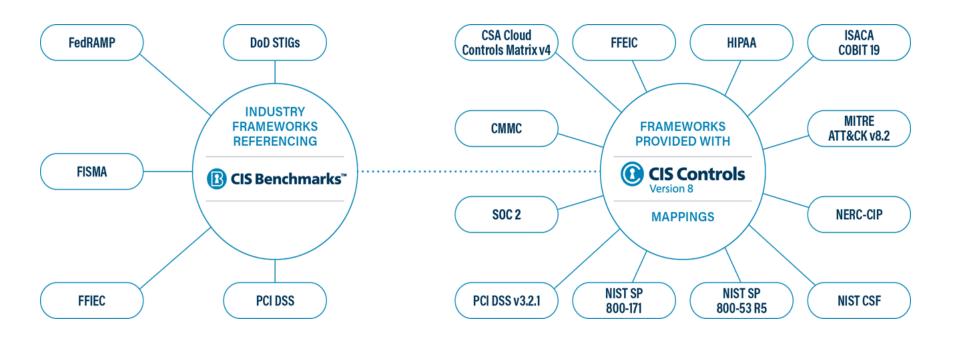


- 100+ CIS Benchmarks covering 25 vendor product families (e.g., AWS Foundations, AWS End User Compute, Amazon EKS)
- Recognized by industry frameworks DISA STIGs, FISMA, FedRAMP, PCI DSS
- Community developed CIS members, SMEs, security community experts, and technology vendors
- Prescriptive instruction
- Step-by-step list to apply configurations
- Rationale on why the configuration is recommended
- Impact the configuration will make
- Mapped to CIS Controls





Scope of Impact of CIS Benchmarks and CIS Controls





Government Configuration Baseline(GCB)



- 政府組態基準(Government Configuration Baseline,以下簡稱GCB)目的在於規範資通訊終端設備(如:個人電腦、伺服器主機及網通設備)的一致性安全設定(如:密碼長度、更新期限等),以降低成為駭客入侵管道,進而引發資安事件之疑慮。
- 國家資通安全研究社(核心業務-資安防護)

1 發展一致的安全組態設定



2 提升政府機關資通訊終端設備使用安全性

GCB 對應法規與落地場域



規定項目:政府組態基準(GCB)

NCCST

資通安全 責任等級 分級辦法

A、B 級公務 機關

初次受核定或等級變更後之一年內,依主管機關公告之項目,完成政府組態基準導入作業,並持續維運。

A、B 級特定 非公務 機關 特定非公務機關之中央目的事業主管機關 得視實際需求,於符合本辦法規定之範圍 內,另行訂定其所管特定非公務機關之資 通安全應辦事項。

CIS Benchmarks 與 GCB 的關係



- 以 Red Hat Enterprise Linux 8
 TWGCB-01-008(v1.0)-政府組態
 基準説明文件。
- 共計 292 項規定,分基本項目 (243 項),防火牆(18 項)以及 SSH 服務(31 項)。
- 主要參考下列二項國際知名資訊 安全標準:
- ➤ CIS (Center for Internet Security)
 RHEL 8 Benchmark
- DISA (Defense Information Systems Agency) RHEL 8 STIG
- > TWGCB for RHEL 5

項次	項目	項數	合計
1	基本項目 (系統設定與維護、網路設定、 日誌與稽核、SELinux)	243	292
2	Firewalld 防火牆組態基準	5	
3	Nftables 防火牆組態基準	7	
4	Iptables 防火牆組態基準	6	
5	SSH 伺服器組態基準	31	





Applying CIS benchmarks in Kubernetes



Covers the classification of Kubernetes



The Types of Kubernetes in CIS Benchmarks is in Server

Software/Kubernetes

Category links:

https://www.cisecurity.org/bench mark/kubernetes Covers the classification of Kubernetes

Red Hat OpenShift Container Platform

Google Kubernetes Engine (GKE)

Amazon Elastic Kubernetes Service (EKS)

Oracle Cloud Infrastructure Container Engine for Kubernetes (OKE)

Azure Kubernetes Service (AKS)

Alibaba Cloud Container Service For Kubernetes (ACK)

Kubernetes

Kubernetes Virtualization

CIS Kubernetes Benchmark v1.9.0

Download PDF

CIS Benchmark Implementation Guide







CIS Kubernetes Benchmark

v1.9.0 - 03-25-2024

Source:

https://www.cisecurity.org/benchmark/kubernetes

1.1.2 Ensure that the API server pod specification file ownership is set to root:root (Automated)

CIS Benchmark措施

Profile Applicability:

· Level 1 - Master Node

Description: 描述

Ensure that the API server pod specification file ownership is set to root:root.

Rationale: 原因理由

The API server pod specification file controls various parameters that set the behavior of the API server. You should set its file ownership to maintain the integrity of the file. The file should be owned by root:root.

Impact: 影響節疇

None

Audit: 如何查核

Run the below command (based on the file location on your system) on the Control Plane node. For example,

stat -c %U:%G /etc/kubernetes/manifests/kube-apiserver.yaml

Verify that the ownership is set to root:root.

Remediation: 建議資訊

Run the below command (based on the file location on your system) on the Control Plane node. For example,

chown root:root /etc/kubernetes/manifests/kube-apiserver.yaml

Default Value: 預設配置

By default, the kube-apiserver.yaml file ownership is set to root:root.

References: 參考資訊

1. https://kubernetes.io/docs/admin/kube-apiserver/

CIS benchmarks Level



Level 1

- More basic
- Smaller service impact
- For small & medium businesses
- For organizations that don't handle PII or other sensitive data

Level 2

- More comprehensive
- Can impact business functions
- For enterprises & critical infrastructure like energy, communications, healthcare, finance, etc.



Applying CIS benchmarks in Kubernetes



1. Control Plane Components

2. Etcd

3. Control Plane Configuration

4. Worker Nodes

5. Policies





Source:

https://www.cisecurity.org/benchmark/kubernetes

1.1 Control Plane Node Configuration Files



- 1.1.1 Ensure that the API server pod specification file permissions are set to 600 or more restrictive (Automated)
- 1.1.2 Ensure that the API server pod specification file ownership is set to root:root (Automated)
- 1.1.3 Ensure that the controller manager pod specification file permissions are set to 600 or more restrictive (Automated)
- 1.1.4 Ensure that the controller manager pod specification file ownership is set to root:root (Automated)
- 1.1.5 Ensure that the scheduler pod specification file permissions are set to 600 or more restrictive (Automated)
- 1.1.6 Ensure that the scheduler pod specification file ownership is set to root:root (Automated)
- 1.1.7 Ensure that the etcd pod specification file permissions are set to 600 or more restrictive (Automated)
- 1.1.8 Ensure that the etcd pod specification file ownership is set to root:root (Automated)
- 1.1.9 Ensure that the Container Network Interface file permissions are set to 600 or more restrictive (Manual)
- 1.1.10 Ensure that the Container Network Interface file ownership is set to root:root (Manual)



1.1 Control Plane Node Configuration Files



- 1.1.11 Ensure that the etcd data directory permissions are set to 700 or more restrictive (Automated)
- 1.1.12 Ensure that the etcd data directory ownership is set to etcd:etcd (Automated)
- 1.1.13 Ensure that the default administrative credential file permissions are set to 600 (Automated)
- 1.1.14 Ensure that the default administrative credential file ownership is set to root:root (Automated)
- 1.1.15 Ensure that the scheduler.conf file permissions are set to 600 or more restrictive (Automated)
- 1.1.16 Ensure that the scheduler.conf file ownership is set to root:root (Automated)
- 1.1.17 Ensure that the controller-manager.conf file permissions are set to 600 or more restrictive (Automated)
- 1.1.18 Ensure that the controller-manager.conf file ownership is set to root:root (Automated)
- 1.1.19 Ensure that the Kubernetes PKI directory and file ownership is set to root:root (Automated)
- 1.1.20 Ensure that the Kubernetes PKI certificate file permissions are set to 600 or more restrictive (Manual)
- 1.1.21 Ensure that the Kubernetes PKI key file permissions are set to 600 (Manual)



1.2 API Server



- 1.2.1 Ensure that the --anonymous-auth argument is set to false (Manual)
- 1.2.2 Ensure that the --token-auth-file parameter is not set (Automated)
- 1.2.3 Ensure that the DenyServiceExternalIPs is set (Manual)
- 1.2.4 Ensure that the --kubelet-client-certificate and --kubeletclient-key arguments are set as appropriate (Automated)
- 1.2.5 Ensure that the --kubelet-certificate-authority argument is set as appropriate (Automated)
- 1.2.6 Ensure that the --authorization-mode argument is not set to AlwaysAllow (Automated)
- 1.2.7 Ensure that the --authorization-mode argument includes Node (Automated)
- 1.2.8 Ensure that the --authorization-mode argument includes RBAC (Automated)
- 1.2.9 Ensure that the admission control plugin EventRateLimit is set (Manual)
- 1.2.10 Ensure that the admission control plugin AlwaysAdmit is not set (Automated)



1.2 API Server



- 1.2.11 Ensure that the admission control plugin AlwaysPullImages is set (Manual)
- 1.2.12 Ensure that the admission control plugin ServiceAccount is set (Automated)
- 1.2.13 Ensure that the admission control plugin NamespaceLifecycle is set (Automated)
- 1.2.14 Ensure that the admission control plugin NodeRestriction is set (Automated)
- 1.2.15 Ensure that the --profiling argument is set to false (Automated)
- 1.2.16 Ensure that the --audit-log-path argument is set (Automated)
- 1.2.17 Ensure that the --audit-log-maxage argument is set to 30 or as appropriate (Automated)
- 1.2.18 Ensure that the --audit-log-maxbackup argument is set to 10 or as appropriate (Automated)
- 1.2.19 Ensure that the --audit-log-maxsize argument is set to 100 or as appropriate (Automated)
- 1.2.20 Ensure that the --request-timeout argument is set as appropriate (Manual)



1.2 API Server



- 1.2.21 Ensure that the --service-account-lookup argument is set to true (Automated)
- 1.2.22 Ensure that the --service-account-key-file argument is set as appropriate (Automated)
- 1.2.23 Ensure that the --etcd-certfile and --etcd-keyfile arguments are set as appropriate (Automated)
- 1.2.24 Ensure that the --tls-cert-file and --tls-private-key-file arguments are set as appropriate (Automated)
- 1.2.25 Ensure that the --client-ca-file argument is set as appropriate (Automated)
- 1.2.26 Ensure that the --etcd-cafile argument is set as appropriate (Automated)
- 1.2.27 Ensure that the --encryption-provider-config argument is set as appropriate (Manual)
- 1.2.28 Ensure that encryption providers are appropriately configured (Manual)
- 1.2.29 Ensure that the API Server only makes use of Strong Cryptographic Ciphers (Manual)



1.3 Controller Manager



- 1.3.1 Ensure that the --terminated-pod-gc-threshold argument is set as appropriate (Manual)
- 1.3.2 Ensure that the --profiling argument is set to false (Automated)
- 1.3.3 Ensure that the --use-service-account-credentials argument is set to true (Automated)
- 1.3.4 Ensure that the --service-account-private-key-file argument is set as appropriate (Automated)
- 1.3.5 Ensure that the --root-ca-file argument is set as appropriate (Automated)
- 1.3.6 Ensure that the RotateKubeletServerCertificate argument is set to true (Automated)
- 1.3.7 Ensure that the --bind-address argument is set to 127.0.0.1 (Automated)



1.4 Scheduler



- 1.4.1 Ensure that the --profiling argument is set to false (Automated)
- 1.4.2 Ensure that the --bind-address argument is set to 127.0.0.1 (Automated)

2. Etcd



- 2.1 Ensure that the --cert-file and --key-file arguments are set as appropriate (Automated)
- 2.2 Ensure that the --client-cert-auth argument is set to true (Automated)
- 2.3 Ensure that the --auto-tls argument is not set to true (Automated)
- 2.4 Ensure that the --peer-cert-file and --peer-key-file arguments are set as appropriate (Automated)
- 2.5 Ensure that the --peer-client-cert-auth argument is set to true (Automated)
- 2.6 Ensure that the --peer-auto-tls argument is not set to true (Automated)
- 2.7 Ensure that a unique Certificate Authority is used for etcd (Manual)

3. Control Plane Configuration



- 3.1 Authentication and Authorization
- 3.1.1 Client certificate authentication should not be used for users (Manual)
- 3.1.2 Service account token authentication should not be used for users (Manual)
- 3.1.3 Bootstrap token authentication should not be used for users (Manual)
- 3.2 Logging
- 3.2.1 Ensure that a minimal audit policy is created (Manual)
- 3.2.2 Ensure that the audit policy covers key security concerns (Manual)



4.1 Worker Node Configuration Files



- 4.1.1 Ensure that the kubelet service file permissions are set to 600 or more restrictive (Automated)
- 4.1.2 Ensure that the kubelet service file ownership is set to root:root (Automated)
- 4.1.3 If proxy kubeconfig file exists ensure permissions are set to 600 or more restrictive (Manual)
- 4.1.4 If proxy kubeconfig file exists ensure ownership is set to root:root (Manual)
- 4.1.5 Ensure that the --kubeconfig kubelet.conf file permissions are set to 600 or more restrictive (Automated)
- 4.1.6 Ensure that the --kubeconfig kubelet.conf file ownership is set to root:root (Automated)
- 4.1.7 Ensure that the certificate authorities file permissions are set to 600 or more restrictive (Manual)



4.2 Kubelet



- 4.2.1 Ensure that the --anonymous-auth argument is set to false (Automated)
- 4.2.2 Ensure that the --authorization-mode argument is not set to AlwaysAllow (Automated)
- 4.2.3 Ensure that the --client-ca-file argument is set as appropriate (Automated)
- 4.2.4 Verify that the --read-only-port argument is set to 0 (Manual)
- 4.2.5 Ensure that the --streaming-connection-idle-timeout argument is not set to 0 (Manual)
- 4.2.6 Ensure that the --make-iptables-util-chains argument is set to true (Automated)
- 4.2.7 Ensure that the --hostname-override argument is not set (Manual)



4.2 Kubelet & 4.3 kube-proxy



- 4.2 Kubelet
- 4.2.8 Ensure that the eventRecordQPS argument is set to a level which ensures appropriate event capture (Manual)
- 4.2.9 Ensure that the --tls-cert-file and --tls-private-key-file arguments are set as appropriate (Manual)
- 4.2.10 Ensure that the --rotate-certificates argument is not set to false (Automated)
- 4.2.11 Verify that the RotateKubeletServerCertificate argument is set to true (Manual)
- 4.2.12 Ensure that the Kubelet only makes use of Strong Cryptographic Ciphers (Manual)
- 4.2.13 Ensure that a limit is set on pod PIDs (Manual)
- 4.3 kube-proxy
- 4.3.1 Ensure that the kube-proxy metrics service is bound to localhost (Automated)



5. Policies



- 5.1 RBAC and Service Accounts
- 5.1.1 Ensure that the cluster-admin role is only used where required (Manual)
- 5.1.2 Minimize access to secrets (Manual)
- 5.1.3 Minimize wildcard use in Roles and ClusterRoles (Manual)
- 5.1.4 Minimize access to create pods (Manual)
- 5.1.5 Ensure that default service accounts are not actively used. (Manual)
- 5.1.6 Ensure that Service Account Tokens are only mounted where necessary (Manual)
- 5.1.7 Avoid use of system:masters group (Manual)
- 5.1.8 Limit use of the Bind, Impersonate and Escalate permissions in the Kubernetes cluster (Manual)
- 5.1.9 Minimize access to create persistent volumes (Manual)
- 5.1.10 Minimize access to the proxy sub-resource of nodes (Manual)
- 5.1.11 Minimize access to the approval sub-resource of certificatesigningrequests objects (Manual)

 CIS Benchmarks™

5. Policies



- 5.2 Pod Security Standards
- 5.2.1 Ensure that the cluster has at least one active policy control mechanism in place (Manual)
- 5.2.2 Minimize the admission of privileged containers (Manual)
- 5.2.3 Minimize the admission of containers wishing to share the host process ID namespace (Manual)
- 5.2.4 Minimize the admission of containers wishing to share the host IPC namespace (Manual)
- 5.2.5 Minimize the admission of containers wishing to share the host network namespace (Manual)
- 5.2.6 Minimize the admission of containers with allowPrivilegeEscalation (Manual)
- 5.2.7 Minimize the admission of root containers (Manual)



Policies



- Pod Security Standards
- 5.2.8 Minimize the admission of containers with the NET_RAW capability (Manual)
- 5.2.9 Minimize the admission of containers with added capabilities (Manual)
- 5.2.10 Minimize the admission of containers with capabilities assigned (Manual)
- 5.2.11 Minimize the admission of Windows HostProcess Containers (Manual)
- 5.2.12 Minimize the admission of HostPath volumes (Manual)
- 5.2.13 Minimize the admission of containers which use HostPorts (Manual)

Policies



- 5.3 Network Policies and CNI
- 5.3.1 Ensure that the CNI in use supports Network Policies (Manual)
- 5.3.2 Ensure that all Namespaces have Network Policies defined (Manual)
- 5.4 Secrets Management
- 5.4.1 Prefer using secrets as files over secrets as environment variables (Manual)
- 5.4.2 Consider external secret storage (Manual)
- 5.5 Extensible Admission Control
- 5.5.1 Configure Image Provenance using ImagePolicyWebhook admission controller (Manual)



5. Policies



- 5.7 General Policies
- 5.7.1 Create administrative boundaries between resources using namespaces (Manual)
- 5.7.2 Ensure that the seccomp profile is set to docker/default in your pod definitions (Manual)
- 5.7.3 Apply Security Context to Your Pods and Containers (Manual)
- 5.7.4 The default namespace should not be used (Manual)





TREND

trivv

trivy

CNCF landscape:Security & Compliance

WhiteSource | Wh

Clarity





Source: https://landscape.cncf.io/



Kube-bench: Kubernetes Security Assessment Too

Overview:

Kube-bench is an open-source tool for assessing the security of Kubernetes clusters. It runs checks based on the Center for Internet Security (CIS) Kubernetes benchmark and was developed in GoLang by Aqua Security.

• Key Features:

- Security Reconnaissance: Supports both self-managed Kubernetes clusters and clusters managed by cloud providers (e.g., AWS, Azure, GKE).
- Figure 1 Test Rule Configuration: Configured via YAML files, allowing easy updates.

CIS Kubernetes Benchmark:

Follows best practices to ensure Kubernetes cluster security.



Reporting and Alerts:

Supports report generation and alert creation based on JSON format scan scripts, with customizable policies.



```
controlplane $ kube-bench run
[INFO] 1 Master Node Security Configuration
[INFO] 1.1 Master Node Configuration Files
[PASS] 1.1.1 Ensure that the API server pod specification file permissions are set to 644 or more restrictive (Automated)
[PASS] 1 1 2 Ensure that the APT server nod specification file ownership is set to root:root (Automated)
   State En CIS Benchmark rule anager pod specification file permissions are set to 644 or more restrictive (Automated)
PASS | 1.1.4 Ensure that the controller manager pod specification file ownership is set to root:root (Automated)
[PASS] 1.1.5 Ensure that the scheduler pod specification file permissions are set to 644 or more restrictive (Automated)
[PASS] 1.1.6 Ensure that the scheduler pod specification file ownership is set to root:root (Automated)
[PASS] 1.1.7 Ensure that the etcd pod specification file permissions are set to 644 or more restrictive (Automated)
[PASS] 1.1.8 Ensure that the etcd pod specification file ownership is set to root:root (Automated)
[WARN] 1.1.9 Ensure that the Container Network Interface file permissions are set to 644 or more restrictive (Manual)
[PASS] 1.1.10 Ensure that the Container Network Interface file ownership is set to root:root (Manual)
[PASS] 1.1.11 Ensure that the etcd data directory permissions are set to 700 or more restrictive (Automated)
      1.1.12 Ensure that the etcd data directory ownership is set to etcd:etcd (Automated)
[PASS] 1.1.13 Ensure that the admin.conf file permissions are set to 644 or more restrictive (Automated)
[PASS] 1.1.14 Ensure that the admin.conf file ownership is set to root:root (Automated)
[PASS] 1.1.15 Ensure that the scheduler.conf file permissions are set to 644 or more restrictive (Automated)
PASS] 1.1.16 Ensure that the scheduler.conf file ownership is set to root:root (Automated)
[PASS] 1.1.17 Ensure that the controller-manager.conf file permissions are set to 644 or more restrictive (Automated)
[PASS] 1.1.18 Ensure that the controller-manager.conf file ownership is set to root:root (Automated)
      1.1.19 Ensure that the Kubernetes PKI directory and file ownership is set to root:root (Automated)
[PASS] 1.1.20 Ensure that the Kubernetes PKI certificate file permissions are set to 644 or more restrictive (Manual)
[PASS]401.1.21 Ensure that the Kubernetes PKI key file permissions are set to 600 (Manual)
```



Туре	Result	Staute
manual		[WARN]
manual	通過	[PASS]
manual	失敗	[FAIL]
automated	通過	[PASS]
automated	失敗	[FAIL]
skip		[INFO]

[PASS]/[FAIL]: Execution success or failure [WARN]: Requires further confirmation [INFO]: No further action needed





```
== Summary policies ==
0 checks PASS
0 checks FAIL
26 checks WARN
0 checks INFO
68 checks PASS
13 checks FAIL
42 checks WARN
0 checks INFO
```

Summary #section

Summary total





```
[INFO] 1.3 Controller Manager
[WARN] 1.3.1 Ensure that the --terminated-pod-gc-threshold argument is set as appropriate (Manual)
      1.3.2 Ensure that the --profiling argument is set to false (Automated)
PASS 13.3 Ensure that the --use-service-account-credentials argument is set to true (Automated)
[PASS] 1.1.4 Ensure that the --service-account-private-key-file argument is set as appropriate (Automated)
      1.3 5 Ensure that the --root-ca-file argument is set as appropriate (Automated)
PASS 1.3. Ensur
                    1.3.1 Edit the Controller Manager pod specification file /etc/kubernetes/manifests/kube-controller-manager.yaml
PASS 1.3.7 Ensur
                    on the master node and set the --terminated-pod-gc-threshold to an appropriate threshold,
INFO] 1.4 Schedul
                    for example:
      1.4.1 E sur
                    --terminated-pod-gc-threshold=10
[PASS] 1.4.2 En ur
                    1.3.2 Edit the Controller Manager pod specification file /etc/kubernetes/manifests/kube-controller-manager.yaml
                    on the master node and set the below parameter.
                    --profiling=false
                    1.4.1 Edit the Scheduler pod specification file /etc/kubernetes/manifests/kube-scheduler.yaml file
                    on the master node and set the below parameter.
                     --profiling=false
```

Remediations Steps



Thanks!



